



Per-layer marginal quality contributions ($\Delta_{k \rightarrow k+1}$, averaged across seven models)

quality gain quality loss

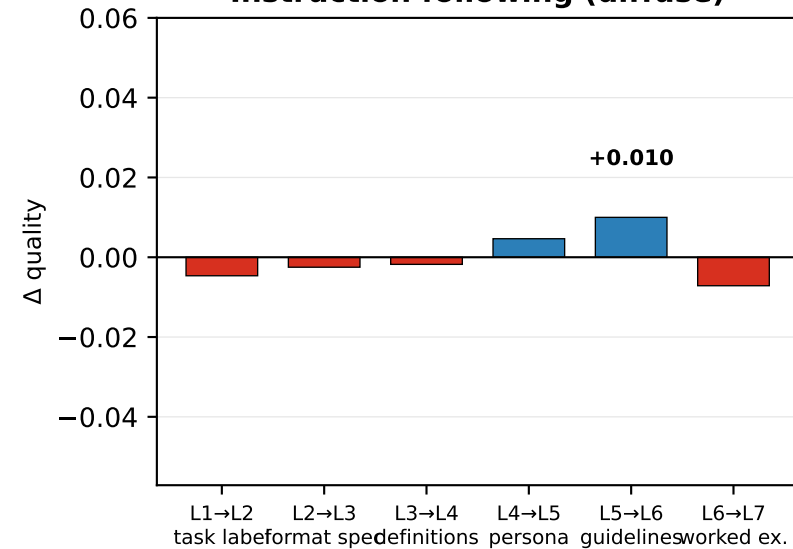
Classification (concentrated)



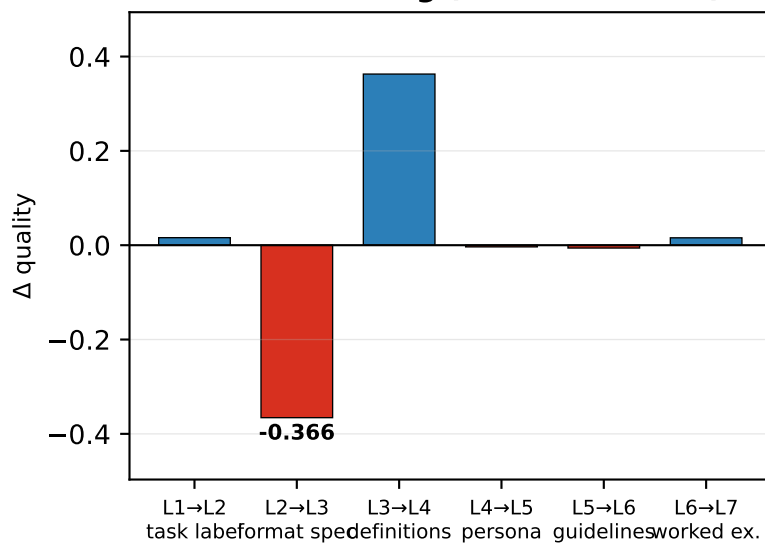
Product extraction (distributed)



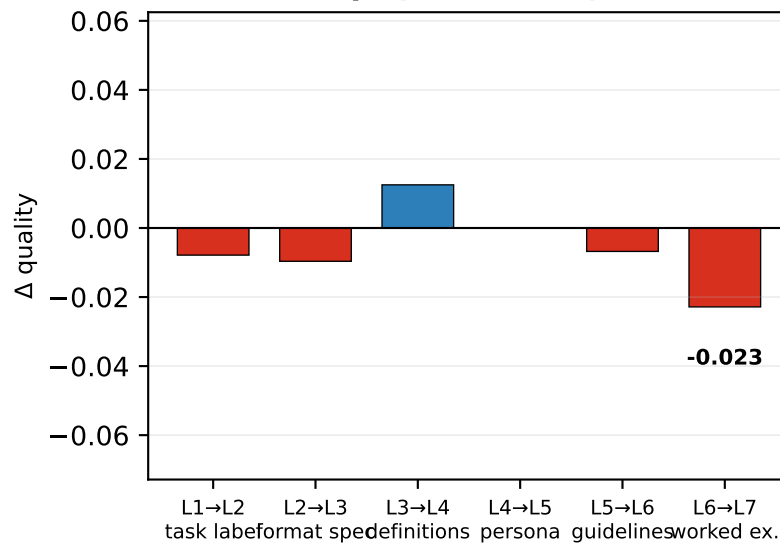
Instruction following (diffuse)



Math reasoning (non-monotonic)



QA (insensitive)



Summarisation (insensitive)

